

```

!pip install datasets &>> install.log

from PIL import Image
import torch.nn.functional as F
import os
from tqdm.notebook import tqdm
import torch
import numpy as np

def img_to_tensor(im):
    return torch.tensor(np.array(im.convert('RGB'))/255).permute(2, 0, 1).unsqueeze(0) * 2 - 1

def tensor_to_image(t):
    return Image.fromarray(np.array(((t.squeeze()).permute(1, 2, 0)+1)/2).clip(0, 1)*255).astype(np.uint8))

def gather(consts: torch.Tensor, t: torch.Tensor):
    """Gather consts for $t$ and reshape to feature map shape"""
    c = consts.gather(-1, t)
    return c.reshape(-1, 1, 1, 1)

import datasets
from datasets import load_dataset

cifar10 = load_dataset('cifar10')

# View some examples:
image = Image.new('RGB', size=(32*5, 32*2))
for i in range(10):
    im = cifar10['train'][i]['img']
    image.paste(im, ( (i%5)*32, (i//5)*32 ))
image.resize((32*5*4, 32*2*4), Image.NEAREST)

/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:89: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your
settings tab (https://huggingface.co/settings/tokens), set it as
secret in your Google Colab and restart your session.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to
access public models or datasets.
  warnings.warn(

{"model_id": "15bc561682344fda9278db858b3f0d02", "version_major": 2, "version_minor": 0}


```

```
{
  "model_id": "470d6f4e18834b3cb26c97ed7722d04b",
  "version_major": 2,
  "version_minor": 0
}

{
  "model_id": "2ffc877d0df5427c8287e6957007d6cd",
  "version_major": 2,
  "version_minor": 0
}

{
  "model_id": "6ad072a8f0bb4b5ab3fcf1b65a79a71d",
  "version_major": 2,
  "version_minor": 0
}

{
  "model_id": "af49226d070f48d989e219bc5be4e718",
  "version_major": 2,
  "version_minor": 0
}
```



##UNet Definition code

```
import math
from typing import Optional, Tuple, Union, List

import torch
from torch import nn

# A fancy activation function
class Swish(nn.Module):
    """
    ### Swish activation function
    $$x \cdot \sigma(x)$$
    """

    def forward(self, x):
        return x * torch.sigmoid(x)

# The time embedding
class TimeEmbedding(nn.Module):
    """
    ### Embeddings for $t$
    """
```

```

def __init__(self, n_channels: int):
    """
    * `n_channels` is the number of dimensions in the embedding
    """
    super().__init__()
    self.n_channels = n_channels
    # First linear layer
    self.lin1 = nn.Linear(self.n_channels // 4, self.n_channels)
    # Activation
    self.act = Swish()
    # Second linear layer
    self.lin2 = nn.Linear(self.n_channels, self.n_channels)

def forward(self, t: torch.Tensor):
    # Create sinusoidal position embeddings
    # [same as those from the
transformer](../../transformers/positional_encoding.html)
    #
    # \begin{align}
    # PE^{(1)}_{t,i} &= \sin\Bigg(\frac{t}{10000^{\frac{i}{d} -
1}}}\Bigg) \setminus \\
    # PE^{(2)}_{t,i} &= \cos\Bigg(\frac{t}{10000^{\frac{i}{d} -
1}}}\Bigg) \\
    # \end{align}
    #
    # where $d$ is `half_dim`
    half_dim = self.n_channels // 8
    emb = math.log(10_000) / (half_dim - 1)
    emb = torch.exp(torch.arange(half_dim, device=t.device) * -
emb)
    emb = t[:, None] * emb[None, :]
    emb = torch.cat((emb.sin(), emb.cos()), dim=1)

    # Transform with the MLP
    emb = self.act(self.lin1(emb))
    emb = self.lin2(emb)

    #
    return emb

# Residual blocks include 'skip' connections
class ResidualBlock(nn.Module):
    """
    ### Residual block
    A residual block has two convolution layers with group
normalization.
    Each resolution is processed with two residual blocks.
    """

```

```

def __init__(self, in_channels: int, out_channels: int,
time_channels: int, n_groups: int = 32):
    """
    * `in_channels` is the number of input channels
    * `out_channels` is the number of input channels
    * `time_channels` is the number channels in the time step
    ($t$) embeddings
    * `n_groups` is the number of groups for [group normalization]
    (.../normalization/group_norm/index.html)
    """
    super().__init__()
    # Group normalization and the first convolution layer
    self.norm1 = nn.GroupNorm(n_groups, in_channels)
    self.act1 = Swish()
    self.conv1 = nn.Conv2d(in_channels, out_channels,
kernel_size=(3, 3), padding=(1, 1))

    # Group normalization and the second convolution layer
    self.norm2 = nn.GroupNorm(n_groups, out_channels)
    self.act2 = Swish()
    self.conv2 = nn.Conv2d(out_channels, out_channels,
kernel_size=(3, 3), padding=(1, 1))

    # If the number of input channels is not equal to the number
    of output channels we have to
    # project the shortcut connection
    if in_channels != out_channels:
        self.shortcut = nn.Conv2d(in_channels, out_channels,
kernel_size=(1, 1))
    else:
        self.shortcut = nn.Identity()

    # Linear layer for time embeddings
    self.time_emb = nn.Linear(time_channels, out_channels)

def forward(self, x: torch.Tensor, t: torch.Tensor):
    """
    * `x` has shape `[batch_size, in_channels, height, width]`
    * `t` has shape `[batch_size, time_channels]`
    """
    # First convolution layer
    h = self.conv1(self.act1(self.norm1(x)))
    # Add time embeddings
    h += self.time_emb(t)[:, :, None, None]
    # Second convolution layer
    h = self.conv2(self.act2(self.norm2(h)))

    # Add the shortcut connection and return
    return h + self.shortcut(x)

```

```

# Ahh yes, magical attention...
class AttentionBlock(nn.Module):
    """
    ### Attention block
    This is similar to [transformer multi-head
    attention](../transformers/mha.html).
    """

    def __init__(self, n_channels: int, n_heads: int = 1, d_k: int =
    None, n_groups: int = 32):
        """
        * `n_channels` is the number of channels in the input
        * `n_heads` is the number of heads in multi-head attention
        * `d_k` is the number of dimensions in each head
        * `n_groups` is the number of groups for [group normalization]
        (../normalization/group_norm/index.html)
        """
        super().__init__()

        # Default `d_k`
        if d_k is None:
            d_k = n_channels
        # Normalization layer
        self.norm = nn.GroupNorm(n_groups, n_channels)
        # Projections for query, key and values
        self.projection = nn.Linear(n_channels, n_heads * d_k * 3)
        # Linear layer for final transformation
        self.output = nn.Linear(n_heads * d_k, n_channels)
        # Scale for dot-product attention
        self.scale = d_k ** -0.5
        #
        self.n_heads = n_heads
        self.d_k = d_k

    def forward(self, x: torch.Tensor, t: Optional[torch.Tensor] =
    None):
        """
        * `x` has shape `[batch_size, in_channels, height, width]`
        * `t` has shape `[batch_size, time_channels]`
        """
        # `t` is not used, but it's kept in the arguments because for
        the attention layer function signature
        # to match with `ResidualBlock`.
        _ = t
        # Get shape
        batch_size, n_channels, height, width = x.shape
        # Change `x` to shape `[batch_size, seq, n_channels]`
        x = x.view(batch_size, n_channels, -1).permute(0, 2, 1)
        # Get query, key, and values (concatenated) and shape it to
        `[batch_size, seq, n_heads, 3 * d_k]`

```

```

        qkv = self.projection(x).view(batch_size, -1, self.n_heads, 3
* self.d_k)
        # Split query, key, and values. Each of them will have shape
        `[batch_size, seq, n_heads, d_k]`
        q, k, v = torch.chunk(qkv, 3, dim=-1)
        # Calculate scaled dot-product  $\frac{Q K^{\top}}{\sqrt{d_k}}$ 
        attn = torch.einsum('bihd,bjhd->bijh', q, k) * self.scale
        # Softmax along the sequence dimension  $\underset{seq}{\text{softmax}}$ 
         $\Bigg(\frac{Q K^{\top}}{\sqrt{d_k}}\Bigg)$ 
        attn = attn.softmax(dim=1)
        # Multiply by values
        res = torch.einsum('bijh,bjhd->bihd', attn, v)
        # Reshape to `[batch_size, seq, n_heads * d_k]`
        res = res.view(batch_size, -1, self.n_heads * self.d_k)
        # Transform to `[batch_size, seq, n_channels]`
        res = self.output(res)

        # Add skip connection
        res += x

        # Change to shape `[batch_size, in_channels, height, width]`
        res = res.permute(0, 2, 1).view(batch_size, n_channels,
height, width)

        #
        return res

```

```

class DownBlock(nn.Module):
    """
    ### Down block
    This combines `ResidualBlock` and `AttentionBlock`. These are used
    in the first half of U-Net at each resolution.
    """

    def __init__(self, in_channels: int, out_channels: int,
time_channels: int, has_attn: bool):
        super().__init__()
        self.res = ResidualBlock(in_channels, out_channels,
time_channels)
        if has_attn:
            self.attn = AttentionBlock(out_channels)
        else:
            self.attn = nn.Identity()

    def forward(self, x: torch.Tensor, t: torch.Tensor):
        x = self.res(x, t)
        x = self.attn(x)
        return x

```

```

class UpBlock(nn.Module):
    """
    ### Up block
    This combines `ResidualBlock` and `AttentionBlock`. These are used
    in the second half of U-Net at each resolution.
    """

    def __init__(self, in_channels: int, out_channels: int,
time_channels: int, has_attn: bool):
        super().__init__()
        # The input has `in_channels + out_channels` because we
        concatenate the output of the same resolution
        # from the first half of the U-Net
        self.res = ResidualBlock(in_channels + out_channels,
out_channels, time_channels)
        if has_attn:
            self.attn = AttentionBlock(out_channels)
        else:
            self.attn = nn.Identity()

    def forward(self, x: torch.Tensor, t: torch.Tensor):
        x = self.res(x, t)
        x = self.attn(x)
        return x


class MiddleBlock(nn.Module):
    """
    ### Middle block
    It combines a `ResidualBlock`, `AttentionBlock`, followed by
    another `ResidualBlock`.
    This block is applied at the lowest resolution of the U-Net.
    """

    def __init__(self, n_channels: int, time_channels: int):
        super().__init__()
        self.res1 = ResidualBlock(n_channels, n_channels,
time_channels)
        self.attn = AttentionBlock(n_channels)
        self.res2 = ResidualBlock(n_channels, n_channels,
time_channels)

    def forward(self, x: torch.Tensor, t: torch.Tensor):
        x = self.res1(x, t)
        x = self.attn(x)
        x = self.res2(x, t)
        return x

```

```

class Upsample(nn.Module):
    """
    ### Scale up the feature map by  $2 \times$ 
    """

    def __init__(self, n_channels):
        super().__init__()
        self.conv = nn.ConvTranspose2d(n_channels, n_channels, (4, 4),
(2, 2), (1, 1))

    def forward(self, x: torch.Tensor, t: torch.Tensor):
        # `t` is not used, but it's kept in the arguments because for
the attention layer function signature
        # to match with `ResidualBlock`.
        _ = t
        return self.conv(x)

class Downsample(nn.Module):
    """
    ### Scale down the feature map by  $\frac{1}{2} \times$ 
    """

    def __init__(self, n_channels):
        super().__init__()
        self.conv = nn.Conv2d(n_channels, n_channels, (3, 3), (2, 2),
(1, 1))

    def forward(self, x: torch.Tensor, t: torch.Tensor):
        # `t` is not used, but it's kept in the arguments because for
the attention layer function signature
        # to match with `ResidualBlock`.
        _ = t
        return self.conv(x)

# The core class definition (aka the important bit)
class UNet(nn.Module):
    """
    ## U-Net
    """

    def __init__(self, image_channels: int = 3, n_channels: int = 64,
ch_mults: Union[Tuple[int, ...], List[int]] = (1, 2,
2, 4),
is_attn: Union[Tuple[bool, ...], List[int]] = (False,
False, True, True),
n_blocks: int = 2):
        """
        * `image_channels` is the number of channels in the image.  $3 \times$ 
for RGB.

```



```

        * `n_channels` is number of channels in the initial feature
map that we transform the image into
        * `ch_mults` is the list of channel numbers at each
resolution. The number of channels is `ch_mults[i] * n_channels`
        * `is_attn` is a list of booleans that indicate whether to use
attention at each resolution
        * `n_blocks` is the number of `UpDownBlocks` at each
resolution
        """
        super().__init__()

        # Number of resolutions
        n_resolutions = len(ch_mults)

        # Project image into feature map
        self.image_proj = nn.Conv2d(image_channels, n_channels,
kernel_size=(3, 3), padding=(1, 1))

        # Time embedding layer. Time embedding has `n_channels * 4`
channels
        self.time_emb = TimeEmbedding(n_channels * 4)

        ##### First half of U-Net - decreasing resolution
        down = []
        # Number of channels
        out_channels = in_channels = n_channels
        # For each resolution
        for i in range(n_resolutions):
            # Number of output channels at this resolution
            out_channels = in_channels * ch_mults[i]
            # Add `n_blocks`
            for _ in range(n_blocks):
                down.append(DownBlock(in_channels, out_channels,
n_channels * 4, is_attn[i]))
                in_channels = out_channels
            # Down sample at all resolutions except the last
            if i < n_resolutions - 1:
                down.append(Downsample(in_channels))

        # Combine the set of modules
        self.down = nn.ModuleList(down)

        # Middle block
        self.middle = MiddleBlock(out_channels, n_channels * 4, )

        ##### Second half of U-Net - increasing resolution
        up = []
        # Number of channels
        in_channels = out_channels
        # For each resolution

```

```

        for i in reversed(range(n_resolutions)):
            # `n_blocks` at the same resolution
            out_channels = in_channels
            for _ in range(n_blocks):
                up.append(UpBlock(in_channels, out_channels,
n_channels * 4, is_attn[i]))
            # Final block to reduce the number of channels
            out_channels = in_channels // ch_mults[i]
            up.append(UpBlock(in_channels, out_channels, n_channels *
4, is_attn[i]))
            in_channels = out_channels
            # Up sample at all resolutions except last
            if i > 0:
                up.append(Upsample(in_channels))

        # Combine the set of modules
        self.up = nn.ModuleList(up)

        # Final normalization and convolution layer
        self.norm = nn.GroupNorm(8, n_channels)
        self.act = Swish()
        self.final = nn.Conv2d(in_channels, image_channels,
kernel_size=(3, 3), padding=(1, 1))

    def forward(self, x: torch.Tensor, t: torch.Tensor):
        """
        * `x` has shape `[batch_size, in_channels, height, width]`
        * `t` has shape `[batch_size]`
        """

        # Get time-step embeddings
        t = self.time_emb(t)

        # Get image projection
        x = self.image_proj(x)

        # `h` will store outputs at each resolution for skip
connection
        h = [x]
        # First half of U-Net
        for m in self.down:
            x = m(x, t)
            h.append(x)

        # Middle (bottom)
        x = self.middle(x, t)

        # Second half of U-Net
        for m in self.up:
            if isinstance(m, Upsample):

```

```

        x = m(x, t)
    else:
        # Get the skip connection from first half of U-Net and
concatenate
        s = h.pop()
        x = torch.cat((x, s), dim=1)
        #
        x = m(x, t)

    # Final normalization and convolution
    return self.final(self.act(self.norm(x)))

```

##Implementation

```

# Create the model
unet = UNet(n_channels=32).cuda()

# Set up some parameters
n_steps = 200
beta = torch.linspace(0.001, 0.1, n_steps).cuda()
alpha = 1. - beta
alpha_bar = torch.cumprod(alpha, dim=0)

# Modified to return the noise itself as well
def q_xt_x0(x0, t):
    mean = gather(alpha_bar, t) ** 0.5 * x0
    var = 1-gather(alpha_bar, t)
    eps = torch.randn_like(x0).to(x0.device)
    return mean + (var ** 0.5) * eps, eps # also returns noise

# Training params
batch_size = 128 # Lower this if hitting memory issues
lr = 1e-3 # Explore this - might want it lower when training on the
full dataset

losses = [] # Store losses for later plotting

dataset = cifar10['train']#.select(range(10000)) # to use a 10k subset
for demo

optim = torch.optim.AdamW(unet.parameters(), lr=lr) # Optimizer

for i in tqdm(range(0, len(dataset)-batch_size, batch_size)): # Run
through the dataset
    ims = [dataset[idx]['img'] for idx in range(i,i+batch_size)] # Fetch
some images
    tims = [img_to_tensor(im).cuda() for im in ims] # Convert to tensors
    x0 = torch.cat(tims) # Combine into a batch
    t = torch.randint(0, n_steps, (batch_size,),
dtype=torch.long).cuda() # Random 't's

```

```


    xt, noise = q_xt_x0(x0, t) # Get the noised images (xt) and the
    noise (our target)
    pred_noise = unet(xt.float(), t) # Run xt through the network to get
    its predictions
    loss = F.mse_loss(noise.float(), pred_noise) # Compare the
    predictions with the targets
    losses.append(loss.item()) # Store the loss for later viewing
    optim.zero_grad() # Zero the gradients
    loss.backward() # Backpropagate the loss (computes and store
    gradients)
    optim.step() # Update the network parameters (using those gradients)

{"model_id": "8c55487b775340cfb144e922360a7806", "version_major": 2, "version_minor": 0}

from matplotlib import pyplot as plt
plt.plot(losses)

[<matplotlib.lines.Line2D at 0x7eec85c23e50>]

```



```

def p_xt(xt, noise, t):
    alpha_t = gather(alpha, t)
    alpha_bar_t = gather(alpha_bar, t)
    eps_coef = (1 - alpha_t) / (1 - alpha_bar_t) ** .5
    mean = 1 / (alpha_t ** 0.5) * (xt - eps_coef * noise) # Note minus

```

```

sign
var = gather(beta, t)
eps = torch.randn(xt.shape, device=xt.device)
return mean + (var ** 0.5) * eps

x = torch.randn(1, 3, 32, 32).cuda() # Start with random noise
ims = []
for i in range(n_steps):
    t = torch.tensor(n_steps-i-1, dtype=torch.long).cuda()
    with torch.no_grad():
        pred_noise = unet(x.float(), t.unsqueeze(0))
        x = p_xt(x, pred_noise, t.unsqueeze(0))
        if i%24 == 0:
            ims.append(tensor_to_image(x.cpu()))

image = Image.new('RGB', size=(32*5, 32))
for i, im in enumerate(ims[:5]):
    image.paste(im, ((i%5)*32, 0))
image.resize((32*4*5, 32*4), Image.NEAREST)

```



Using this set of hyperparameters produces terrible results, even though the loss curve comes similar, due to the higher value of noising schedule(beta) the model can't effectively work.

Using different sets of hyperparameters

```

# Create the model
unet = UNet(n_channels=32).cuda()

# Set up some parameters
n_steps = 100
beta = torch.linspace(0.0001, 0.04, n_steps).cuda()
alpha = 1. - beta
alpha_bar = torch.cumprod(alpha, dim=0)

# Modified to return the noise itself as well
def q_xt_x0(x0, t):
    mean = gather(alpha_bar, t) ** 0.5 * x0
    var = 1-gather(alpha_bar, t)
    eps = torch.randn_like(x0).to(x0.device)
    return mean + (var ** 0.5) * eps, eps # also returns noise

# Training params

```

```

batch_size = 128 # Lower this if hitting memory issues
lr = 2e-4 # Explore this - might want it lower when training on the
full dataset

losses = [] # Store losses for later plotting

dataset = cifar10['train'].select(range(10000)) # to use a 10k subset
for demo

optim = torch.optim.AdamW(unet.parameters(), lr=lr) # Optimizer


for i in tqdm(range(0, len(dataset)-batch_size, batch_size)): # Run
through the dataset
    ims = [dataset[idx]['img'] for idx in range(i,i+batch_size)] # Fetch
some images
    tims = [img_to_tensor(im).cuda() for im in ims] # Convert to tensors
    x0 = torch.cat(tims) # Combine into a batch
    t = torch.randint(0, n_steps, (batch_size,),
dtype=torch.long).cuda() # Random 't's
    xt, noise = q_xt_x0(x0, t) # Get the noised images (xt) and the
noise (our target)
    pred_noise = unet(xt.float(), t) # Run xt through the network to get
its predictions
    loss = F.mse_loss(noise.float(), pred_noise) # Compare the
predictions with the targets
    losses.append(loss.item()) # Store the loss for later viewing
    optim.zero_grad() # Zero the gradients
    loss.backward() # Backpropagate the loss (computes and store
gradients)
    optim.step() # Update the network parameters (using those gradients)

{"model_id": "fcde8c917614434c911bd0f61ea4fe07", "version_major": 2, "vers
ion_minor": 0}

plt.plot(losses)

[<matplotlib.lines.Line2D at 0x7eec846ab8b0>]

```




Generating images with this model

```
def p_xt(xt, noise, t):
    alpha_t = gather(alpha, t)
    alpha_bar_t = gather(alpha_bar, t)
    eps_coef = (1 - alpha_t) / (1 - alpha_bar_t) ** .5
    mean = 1 / (alpha_t ** 0.5) * (xt - eps_coef * noise) # Note minus
sign
    var = gather(beta, t)
    eps = torch.randn(xt.shape, device=xt.device)
    return mean + (var ** 0.5) * eps

x = torch.randn(1, 3, 32, 32).cuda() # Start with random noise
ims = []
for i in range(n_steps):
    t = torch.tensor(n_steps-i-1, dtype=torch.long).cuda()
    with torch.no_grad():
        pred_noise = unet(x.float(), t.unsqueeze(0))
        x = p_xt(x, pred_noise, t.unsqueeze(0))
        if i%24 == 0:
            ims.append(tensor_to_image(x.cpu()))


image = Image.new('RGB', size=(32*5, 32))
for i, im in enumerate(ims[:5]):
    image.paste(im, ((i%5)*32, 0))
image.resize((32*4*5, 32*4), Image.NEAREST)
```



```
x = torch.randn(10, 3, 32, 32).cuda() # Start with random noise
ims = []
for i in range(n_steps):
    t = torch.tensor(n_steps-i-1, dtype=torch.long).cuda()
    with torch.no_grad():
        pred_noise = unet(x.float(), t.unsqueeze(0))
        x = p_xt(x, pred_noise, t.unsqueeze(0))

for i in range(10):
    ims.append(tensor_to_image(x[i].unsqueeze(0).cpu()))

image = Image.new('RGB', size=(32*5, 32*2))
for i, im in enumerate(ims):
    image.paste(im, ((i%5)*32, 32*(i//5)))
image.resize((32*4*5, 32*4*2), Image.NEAREST)
```



##Starting with a partially noised sample

from 50 iterations

```
bird = cifar10['train'][7]['img']
x0 = img_to_tensor(bird)
x = torch.cat([q_xt_x0(x0.cuda()), torch.tensor(50,
dtype=torch.long).cuda())[0] for _ in range(10)] )
example_start = q_xt_x0(x0.cuda()), torch.tensor(50,
dtype=torch.long).cuda())[0]
```




```

print(x.shape)
ims = []
for i in range(50, n_steps):
    t = torch.tensor(n_steps-i-1, dtype=torch.long).cuda()
    with torch.no_grad():
        pred_noise = unet(x.float(), t.unsqueeze(0))
        x = p_xt(x, pred_noise, t.unsqueeze(0))

for i in range(10):
    ims.append(tensor_to_image(x[i].unsqueeze(0).cpu()))

image = Image.new('RGB', size=(32*5, 32*2))
for i, im in enumerate(ims):
    image.paste(im, ((i%5)*32, 32*(i//5)))
    if
i==0:image.paste(tensor_to_image(example_start.unsqueeze(0).cpu()),
((i%5)*32, 32*(i//5))) # Show the heavily noised starting point top
left
image.resize((32*4*5, 32*4*2), Image.NEAREST)
torch.Size([10, 3, 32, 32])

```



from 10 iterations

```

bird = cifar10['train'][7]['img']
x0 = img_to_tensor(bird)
x = torch.cat([q_xt_x0(x0.cuda(), torch.tensor(10,
dtype=torch.long).cuda())[0] for _ in range(10)] )
example_start = q_xt_x0(x0.cuda(), torch.tensor(10,
dtype=torch.long).cuda())[0]
print(x.shape)
ims = []

for i in range(50, n_steps):

```

```

t = torch.tensor(n_steps-i-1, dtype=torch.long).cuda()
with torch.no_grad():
    pred_noise = unet(x.float(), t.unsqueeze(0))
    x = p_xt(x, pred_noise, t.unsqueeze(0))

for i in range(10):
    ims.append(tensor_to_image(x[i].unsqueeze(0).cpu()))

image = Image.new('RGB', size=(32*5, 32*2))
for i, im in enumerate(ims):
    image.paste(im, ((i%5)*32, 32*(i//5)))
    if
i==0:image.paste(tensor_to_image(example_start.unsqueeze(0).cpu()),
((i%5)*32, 32*(i//5))) # Show the heavily noised starting point top
left
image.resize((32*4*5, 32*4*2), Image.NEAREST)

torch.Size([10, 3, 32, 32])

```



from 5 iterations

```

bird = cifar10['train'][7]['img']
x0 = img_to_tensor(bird)
x = torch.cat([q_xt_x0(x0.cuda()), torch.tensor(5,
dtype=torch.long).cuda())[0] for _ in range(10)] )
example_start = q_xt_x0(x0.cuda()), torch.tensor(5,
dtype=torch.long).cuda())[0]
print(x.shape)
ims = []
for i in range(50, n_steps):
    t = torch.tensor(n_steps-i-1, dtype=torch.long).cuda()
    with torch.no_grad():
        pred_noise = unet(x.float(), t.unsqueeze(0))

```

```

x = p_xt(x, pred_noise, t.unsqueeze(0))
for i in range(10):
    ims.append(tensor_to_image(x[i].unsqueeze(0).cpu()))
image = Image.new('RGB', size=(32*5, 32*2))
for i, im in enumerate(ims):
    image.paste(im, ((i%5)*32, 32*(i//5)))
    if
i==0:image.paste(tensor_to_image(example_start.unsqueeze(0).cpu()),
((i%5)*32, 32*(i//5))) # Show the heavily noised starting point top
left
image.resize((32*4*5, 32*4*2), Image.NEAREST)
torch.Size([10, 3, 32, 32])

```

